

Problem Sheet 3

```

• begin
•     using Pkg
•     Pkg.activate(".")
•     Pkg.add(["Distributions", "Plots", "PlutoUI"])
•     using Distributions
•     using LinearAlgebra
•     using Plots
•     using PlutoUI
•     default(linewidth = 3.0, legendfontsize= 15.0)
• end

```

Table of Contents

Problem Sheet 3

1. Bayes inference for the variance of a Gaussian
 - (a) [MATH] Show that the posterior probability $p(\lambda|D)$ of the inverse variance is also a Gamma distribution.
 - (b) [MATH] Compute the mean of the posterior distribution of λ . Compare the result with the result ...
 - (c) [MATH] Show that the variance of the posterior distribution shrinks to zero as $N \rightarrow \infty$. Here we have used the notation $\langle \cdot \rangle$ for posterior expectations.
 - (d) [MATH] Show that the predictive distribution is where α_p and β_p were defined above. Note, this is not a Gaussian!
 - (e) [CODE] Program a function generating samples from a known value λ and compute both the pos...
2. Hyperparameter estimation for a generalised linear model
 - (a) [MATH] The posterior distribution $p(w|D)$ of the vector of weights is a Gaussian. Compute the po...
 - (b) [MATH] Derive an EM algorithm for optimising the hyperparameter β by maximising the log-evi...
 - (c) [CODE] Implement the posterior solution of a generalised linear model given an arbitrary base f...
 - (d) [CODE] Given some random data, implement the EM algorithm to optimize α and β

1. Bayes inference for the variance of a Gaussian

Use a Bayesian approach to estimate the inverse variance λ of a univariate Gaussian distribution

$$p(x|\lambda) = \sqrt{\frac{\lambda}{2\pi}} \exp\left[-\frac{\lambda x^2}{2}\right].$$

Here we have assumed for simplicity that the data has zero mean $\mu = 0$. To apply Bayesian inference we specify a **Gamma** prior distribution for λ ,

$$p(\lambda) = \frac{\lambda^{\alpha-1} \exp[-\lambda/\beta]}{\Gamma(\alpha)\beta^\alpha}$$

where the positive numbers α and β , the *hyperparameters* of the model are assumed to be known and $\Gamma(\alpha)$ is Euler's **gamma** function. We then observe a dataset $D = (x_1, x_2, \dots, x_N)$ comprising N independent random samples from $p(x|\lambda)$.

(a) [MATH] Show that the posterior probability $p(\lambda|D)$ of the inverse variance is also a Gamma distribution with parameters

$$\alpha_p = \alpha + \frac{N}{2}, \quad \frac{1}{\beta_p} = \frac{1}{\beta} + \frac{1}{2} \sum_{i=1}^N x_i^2.$$

Write your answer here or on paper

(b) [MATH] Compute the mean of the posterior distribution of λ . Compare the result with the result from the maximum-likelihood estimation, $\lambda_{\text{ML}} = 1/\sigma_{\text{ML}}^2$ and explain what happens if $N \rightarrow \infty$.

Write your answer here or on paper

(c) [MATH] Show that the variance of the posterior distribution

$$V[\lambda_{\text{post}}] = \langle \lambda^2 \rangle - \langle \lambda \rangle^2$$

shrinks to zero as $N \rightarrow \infty$. Here we have used the notation $\langle \cdot \rangle$ for posterior expectations.

Write your answer here or on paper

(d) [MATH] Show that the predictive distribution is

$$p(x|D) = \frac{1}{\sqrt{2\pi}} \frac{\Gamma(\alpha_p + 1/2)}{\Gamma(\alpha_p)} \sqrt{\beta_p} \left(1 + \frac{x^2 \beta_p}{2}\right)^{-\alpha_p - 1/2}$$

where α_p and β_p were defined above. Note, this is not a Gaussian!

Write your answer here or on paper

(e) [CODE] Program a function generating samples from a known value λ and compute both the posterior distribution and ML estimator by adding progressively new samples.

```
• gamma_params = (α=2.0, β=0.5);
```

```
true_λ = 1.127053974465347
```

```
• true_λ = rand(Gamma(gamma_params...)) # We sample a random value λ
```

```
syntax: invalid syntax (incomplete #<julia: "incomplete: premature end of input">)
```

```
1. top-level scope @ none:1
```

```
• generate_x(λ) = ## !! FILL IN !! Generate a random value x from a give λ
```

```
syntax: invalid syntax (incomplete #<julia: "incomplete: premature end of input">)
```

```
1. top-level scope @ none:1
```

```
• αp(x, α) = ## !! FILL IN !! Return the posterior alpha paramter
```

```
syntax: invalid syntax (incomplete #<julia: "incomplete: premature end of input">)
```

```
1. top-level scope @ none:1
```

```
• βp(x, β) = ## !! FILL IN !! Return the posterior beta paramter
```

```
• posterior_λ(x, α, β) = Gamma(αp(x, α), βp(x, β)); # Posterior distribution
```

```
syntax: invalid syntax (incomplete #<julia: "incomplete: premature end of input">)
```

```
1. top-level scope @ none:1
```

```
• λ_ML(x) = ## !! FILL IN !! Return the maximum likelihood estimator
```

0.0:0.016722408026755852:5.0

```

• begin # Plotting values
•     N = 10000
•     bins = range(-5, 5, length = 50)
•     range_λ = range(0, 5.0, length = 300)
• end

```

UndefVarError: generate_x not defined

1. top-level scope @ **Local: 5**

```

• begin
•     xs = Float64[]
•     anim = Animation()
•     for i in 1:N
•         push!(xs, generate_x(true_λ))
•         if i % 10^floor(Int64, log10(i)) == 0
•             p_x = histogram(xs, bins = bins, lw = 0.0, label = "", title = "N = $i")
•             p_λ = plot(range_λ, x -> pdf(posterior_λ(xs, gamma_params...), x), label =
"p(λ|D)")
•             vline!([λ_ML(xs)], label = "λ ML")
•             vline!([true_λ], label = "true λ")
•             plot(p_x, p_λ, size = (800,400))
•             frame(anim)
•         end
•     end
• end

```

ArgumentError: Cannot build empty animations

```

1. var"#buildanimation#214"(::Int64, ::Int64, ::Bool, ::Bool, ::Bool,
::typeof(Plots.buildanimation), ::Plots.Animation, ::String,
::Bool) @ animation.jl:79
2. #gif#210 @ animation.jl:64 [inlined]
3. top-level scope @ Local: 1 [inlined]

```

```

• gif(anim, fps = 6)

```

UndefVarError: ap not defined

```

1. posterior_λ(::Vector{Float64}, ::Float64, ::Float64) @ Other: 1
2. top-level scope @ Local: 1

```

```

• posterior_λ(xs, gamma_params...)

```

2. Hyperparameter estimation for a generalised linear model

Consider a model for a set of data $D = (y_1, \dots, y_n)$ defined by

$$p(D|\mathbf{w}, \beta) = \left(\frac{\beta}{2\pi}\right)^{N/2} \exp \left[-\sum_{i=1}^N \frac{\beta}{2} \left(y_i - \sum_{j=1}^K w_j \Phi_j(x_i) \right)^2 \right]$$

with a fixed set $\{\Phi_1(x), \dots, \Phi_K(x)\}$ of K basis functions. The prior distribution on the weights is given by

$$p(\mathbf{w}|\alpha) = \left(\frac{\alpha}{2\pi}\right)^{K/2} \exp\left[-\frac{\alpha}{2} \sum_{j=1}^K w_j^2\right].$$

This **generalised linear model** assumes that the observations are generated from a weighted linear combination of the basis functions with additive Gaussian noise.

- **(a) [MATH] The posterior distribution $p(\mathbf{w}|D)$ of the vector of weights is a Gaussian. Compute the posterior mean vector $E[\mathbf{w}]$ and the posterior covariance in terms of the matrix \mathbf{X} where $X_{lk} = \Phi_k(x_l)$.**

Write your answer here or on paper

- **(b) [MATH] Derive an EM algorithm for optimising the hyperparameter β by maximising the log-evidence**

$$p(D|\alpha, \beta) = \int p(D|\mathbf{w}, \beta) p(\mathbf{w}|\alpha) d\mathbf{w}$$

Tip

Hint: Treat the weights \mathbf{w} as a set of latent variables similar to the procedure for α given in the lecture. Express your result in terms of the posterior mean and variance.

Write your answer here or on paper

- **(c) [CODE] Implement the posterior solution of a generalised linear model given an arbitrary base function $\Phi(X) = \{\Phi_1(X), \dots, \Phi_K(X)\}$**

Φ (generic function with 2 methods)

```

• begin
•   #  $\Phi(x::Real) = [x]$  # Linear Case
•   #  $\Phi(x::Real) = [1, x, \sin(x), \cos(x)]$ 
•    $\Phi(x::Real) = [\text{one}(x), x, x^2, x^3]$  #evalpoly(x, ones(4))
•    $\Phi(x::Vector) = \text{mapreduce}(\Phi, \text{hcat}, x)$  # Create a matrix out of a vector
• end

```

syntax: invalid identifier name "..."

1. top-level scope @ none:1

```

• function posterior_params(val_ $\Phi$ , y,  $\alpha$ ,  $\beta$ )
•   ## !! FILL IN !!
•   ## This function should return the mean and the covariance of the posterior
•    $\mu = \dots$ 
•    $\Sigma = \dots$ 
•   return  $\mu$ ,  $\Sigma$ 
• end

```

(d) [CODE] Given some random data, implement the EM algorithm to optimize α and β


data points :  10

$\beta =$  1.0

```

• begin # Create some data
•   X = rand(Uniform(-5, 5), Nx); # Sample X uniformly
•   offset = 3
•   X_test = collect(range(-5 - offset, 5 + offset, length = 500))
•   w_true = [2.0, -0.1, 0.5, -0.2]
•   f = evalpoly(X, Ref(w_true)); # return w[0] + w[1] X + w[2] X^2 + ....
•   y = f + randn(Nx) / sqrt( $\beta_{\text{true}}$ ); # Add some noise
• end;

```



syntax: invalid syntax (incomplete #<julia: "incomplete: premature end of input">)

1. top-level scope @ none:1

```
• update_α(μ, Σ, K) = ## !! FILL IN !! Return the optimal alpha parameter
```

syntax: invalid syntax (incomplete #<julia: "incomplete: premature end of input">)

1. top-level scope @ none:1

```
• update_β(val_Φ, y, μ, Σ) = ## !! FILL IN !! Return the optimal beta parameter
```

expectation_step (generic function with 1 method)

```
• function expectation_step(val_Φ, y, μ, Σ, K, α, β)
•     ## !! FILL IN !! Return the value of L
• end
```

UndefVarError: posterior_params not defined

1. top-level scope @ [Local: 8 [inlined]]
2. top-level scope @ none:0

```
• begin
•     α = 1.0 # Initial parameters
•     β = 1.0
•     val_Φ = Φ(X) # Feature map
•     T = 10 # Number of steps
•     K = size(val_Φ, 1) # Feature map dimension
•     for i in 1:T
•         μ, Σ = posterior_params(val_Φ, y, α, β)
•         println("i = $i, pre L = $(expectation_step(val_Φ, y, μ, Σ, K, α, β)), α = $α,
β = $β")
•         α = update_α(μ, Σ, K)
•         β = update_β(val_Φ, y, μ, Σ)
```

```

• println("i = $i, post L = $(expectation_step(val_Φ, y, μ, Σ, K, α, β)), α =
  $α, β = $β")
• end
• end

```

UndefVarError: posterior_params not defined

1. top-level scope @ **Local: 2**

```

• begin
•   μ, Σ = posterior_params(val_Φ, y, α, β)
•   plot(
•     bar([μ, w_true], label = ["E[w]" "True w"], alpha = 0.5, lw = 0.0),
•     heatmap(Σ, title = "Cov(w)", yflip = true),
•   )
• end

```

UndefVarError: μ not defined

1. top-level scope @ **Local: 3**

```

• begin
•   scatter(X, y, lab= "y", xlabel="x", ylabel="y")
•   plot!(X_test, Φ(X_test)' * μ; lab="Prediction", linewidth=5.0)
•   plot!(X_test, evalpoly.(X_test, Ref(w_true)); linestyle=:dash, color=:black,
•     label="f")
• end

```

We can also sample from the posterior to visualize all the different possibilities

 10

UndefVarError: Σ not defined

1. top-level scope @ **Local: 5** [inlined]
 2. top-level scope @ *none:0*

```

• begin
•   p = scatter(X, y, lab= "y", xlabel="x", ylabel="y")
•   S = 100
•   for i in 1:S
•     w = rand(MvNormal(μ, Symmetric(Σ)))
•     plot!(X_test, Φ(X_test)' * w; lab="", color=:black, alpha=0.01)
•   end
•   p
• end

```

(α = 1.0, β = 1.0)

• (;α, β)

1.001

• β_true

fill_in (generic function with 1 method)

