

# Problem Sheet 4

---

```
• begin
•   using Distributions
•   using PlutoUI
•   using LinearAlgebra
•   using DelimitedFiles
•   using DataFrames
•   using KernelFunctions
•   using Plots
•   using StatsPlots
•   default(;linewidth=3.0, legendfontsize=15.0)
• end
```

## 1. Gaussian process (GP) regression

---

For the GP regression problem, we assume that data are generated as

$$y_i = f(x_i) + \nu_i \quad i = 1, \dots, n$$

where the  $\nu_i$  are independent, zero mean Gaussian noise variables within  $E[\nu_i^2] = \sigma^2$  and  $f(\cdot)$  has a GP prior with kernel  $K(x, x')$ .

**(a) [MATH] Show that the Bayesian evidence is given by**

$$p(\mathbf{y}) = \frac{1}{(2\pi)^{n/2} |\det(\mathbf{K} + \sigma^2 \mathbf{I})|^{\frac{1}{2}}} \exp \left[ -\frac{1}{2} \mathbf{y}^T (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} \right]$$

**where  $\mathbf{y} = (y_1, \dots, y_n)$  and the kernel matrix is defined by  $\mathbf{K}_{ij} = K(x_i, x_j)$ .**

### Tip

Calculate the joint density of  $\mathbf{y}$  and use the fact that  $f(x_j)$  and  $\nu_i$  are independent Gaussian random variables. Hence you can add the respective covariance matrices.

*Write your answer here or on paper*

## (b) [CODE] Create a Gaussian prior with zero mean and a RBF Kernel and sample from it on a grid

```
• begin
•   x_test = range(0, 10, length = 200)
•   k = SqExponentialKernel() # This computes  $k(x, x') = \exp(-0.5||x - x'||^2)$ 
•   K = kernelmatrix(k, x_test) + 1e-5I
•   prior_f = ## !! FILL IN !! Give the prior distribution
•   S = 10 # Number of GP samples
• end;
```

Expects 200 elements in each col of y, found 10.

```
1. error(::String) @ error.jl:33
2. _compute_xyz(::StepRangeLen{Float64, Base.TwicePrecision{Float64},
Base.TwicePrecision{Float64}}, ::Vector{Float64}, ::Nothing,
::Bool) @ series.jl:97
3. macro expansion @ series.jl:163 [inlined]
4. apply_recipe(::AbstractDict{Symbol, Any}, ::Type{RecipesPipeline.SliceIt}, ::Any,
::Any, ::Any) @ RecipesBase.jl:282
5. _process_userrecipes! (::Any, ::Any, ::Any) @ user_recipe.jl:36
6. recipe_pipeline! (::Any, ::Any, ::Any) @ RecipesPipeline.jl:70
7. _plot! (::Plots.Plot, ::Any, ::Any) @ plot.jl:172
8. #plot#148 @ plot.jl:58 [inlined]
9. top-level scope @ [Local: 1] [inlined]
```

```
• plot(x_test, rand(prior_f, S); xlabel="x", ylabel="f", label="", alpha=0.5)
```

## (c) [MATH] Given a set of training data $(X, y)$ , compute the predictive distribution of some test data $X_{\text{test}}$

Write your answer here or on paper

## (d) [CODE] Implement the predictive distribution and plot the predictive mean along with one standard error

N  20

log $\sigma$   -1

```
• begin
•    $\sigma = \exp(\text{log}\sigma)$ 
•   X = rand(Uniform(0, 10), N)
•   y = sin.(X) + randn(N) *  $\sigma$ 
• end;
```

Cyclic references among m, pred\_mean\_and\_cov C

```
• function pred_mean_and_cov(k, x_test, x, y)
•   Kx = kernelmatrix(k, x)
```

```

• Kxtest_x = kernelmatrix(k, x_test, x)
• Kxtest = kernelmatrix(k, x_test) + 1e-5I
• ## !! FILL IN !! This should return the mean and the covariance of the
  predictions
• return m, C
• end;

```

Cyclic references among m, pred\_mean\_and\_cov C

```
• m, C = pred_mean_and_cov(k, x_test, X, y);
```

UndefinedVarError: C not defined

1. top-level scope @ ( Local: 2

```

• begin
• plot(x_test, rand(MvNormal(m, Symmetric(C)), S * 10), color=:black, label="",
  alpha=0.1)
• plot!(x_test, m, ribbon = sqrt.(diag(C)), color=:blue, fillalpha=0.2, label =
  "Prediction")
• plot!(x_test, sin.(x_test), color=:red, label="f")
• scatter!(X, y, color=:green, label = "Data")
• end

```

## 2. Gibbs sampler for outlier detection

The file *outlier.dat* on the web page of the course contains a data set  $D = (y_1, \dots, y_N)$ . Most of the observations have been drawn from a Gaussian probability distribution  $\mathcal{N}(y_i; \mu, \sigma^2)$  with mean  $\mu$  and variance  $\sigma^2$ . However,  $D$  contains some **outliers**, which occur with probability  $\epsilon$  and are displaced by a random offset  $A_i$ . For the purpose of **outlier detection** the model is augmented with an indicator variable

$$\delta_i = \begin{cases} 1 & \text{if } y_i \text{ is an outlier,} \\ 0 & \text{if } y_i \text{ is a normal data point,} \end{cases}$$

for each observation. Assuming conjugate priors for the parameters yields the full stochastic model

$$\begin{aligned} \mu &\sim \mathcal{N}(\theta, v^2), & \sigma^{-2} &\sim \text{Gamma}(\kappa, \lambda), & \epsilon &\sim \text{Beta}(\alpha, \beta), \\ y_i &\sim \mathcal{N}(\mu + \delta_i A_i, \sigma^2), & \delta_i &\sim \text{Bernoulli}(\epsilon), & A_i &\sim \mathcal{N}(0, \tau^2). \end{aligned}$$

We want to use a Gibbs sampler in order to draw samples from the posterior  $p(\mu, \sigma^2, \epsilon, \boldsymbol{\delta}, \mathbf{A} | D)$  with  $\boldsymbol{\delta} = (\delta_1, \dots, \delta_N)$  and  $\mathbf{A} = (A_1, \dots, A_N)$ . Some conditional posteriors are given by

$$\begin{aligned} \mu &\sim \mathcal{N}\left(\frac{\sigma^2 \theta + v^2 \sum_{i=1}^N (y_i - \delta_i A_i)}{\sigma^2 + N v^2}, \frac{\sigma^2 v^2}{\sigma^2 + N v^2}\right), \\ \sigma^{-2} &\sim \text{Gamma}\left(\kappa + \frac{N}{2}, \frac{2\lambda}{2 + \lambda \sum_{i=1}^N (y_i - \delta_i A_i - \mu)^2}\right). \end{aligned}$$

- (a) [MATH] Show that the remaining conditional posteriors are given by

$$\begin{aligned}\delta_i &\sim \text{Bernoulli}\left(\frac{\epsilon}{\epsilon + (1 - \epsilon) \exp(-A_i(y_i - A_i - \mu)/(2\sigma^2))}\right), \\ A_i &\sim \mathcal{N}\left(\frac{\tau^2 \delta_i (y_i - \mu)}{\sigma^2 + \tau^2}, \frac{\sigma^2 \tau^2}{\sigma^2 + \tau^2 \delta_i}\right), \\ \epsilon &\sim \text{Beta}\left(\alpha + \sum_{i=1}^N \delta_i, \beta + \sum_{i=1}^N (1 - \delta_i)\right).\end{aligned}$$

Write your answer here or on paper

- (b) [CODE] Write a program that implements the *Gibbs sampler*. Generate  $10^3$  samples from the posterior using the hyperparameters  $\theta = 0$ ,  $v^2 = 100$ ,  $\kappa = 2$ ,  $\lambda = 2$ ,  $\alpha = 2$ ,  $\beta = 20$ ,  $\tau^2 = 100$ . Plot histograms showing the marginal posteriors  $p(\mu|D)$  and  $p(\epsilon|D)$ .

## Solution

```
• function sample_mu(σ², θ, v², y, δ, A, N)
•     ## !! FILL IN !! it should return a sample for μ
•     return
• end;
```

```
• function sample_σ²(κ, N, λ, y, δ, A, μ)
•     ## !! FILL IN !! it should return a sample for σ²
•     return
• end;
```

```
• function sample_δ(ε, A, y, μ, σ²)
•     ## !! FILL IN !! it should return a vector of samples for δ
•     return
• end;
```

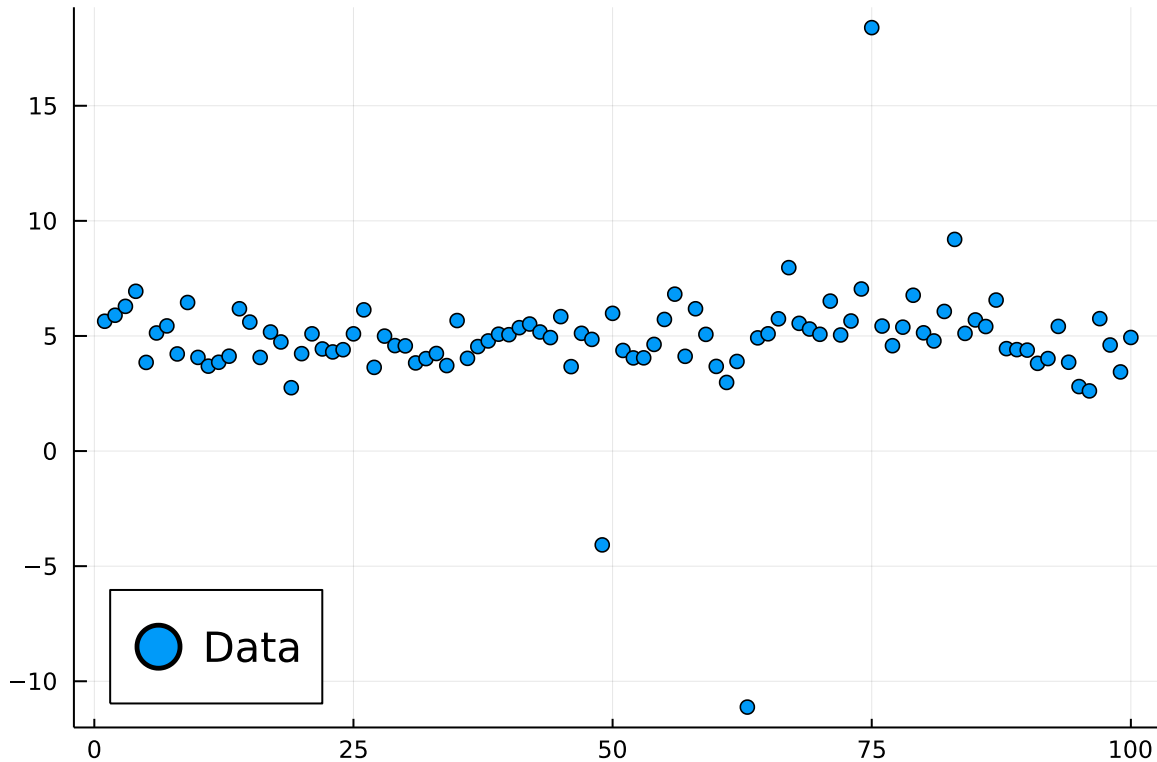
```
• function sample_A(τ², δ, y, μ, σ²)
•     ## !! FILL IN !! it should return a vector of samples for A
•     return
• end;
```

```
• function sample_ε(α, δ, β)
•     ## !! FILL IN !! it should return a sample for \epsilon
•     return
```

```
• end;
```

100

```
• begin # We load the data
•   y_outlier = vec(readlm("outlier.dat"))
•   Ny = length(y_outlier)
• end
```



```
• begin # We select multiple hyperparameters
•   T = 10000
•   θ = 0.0
•   ν² = 100
•   κ = 2
•   λ = 2
•   α = 2
•   β = 20
•   τ² = 100
• end;
```

**MethodError: Cannot `convert` an object of type Nothing to an object of type Float64**

Closest candidates are:

`convert(::Type{S}, !Matched::CategoricalArrays.CategoricalValue)` where `S<:Union{AbstractChar, AbstractString, Number}` at `/home/theo/.julia/packages/CategoricalArrays/Fr04b/src/value.jl:79`

`convert(::Type{T}, !Matched::Base.TwicePrecision)` where `T<:Number` at `twiceprecision.jl:250`

`convert(::Type{T}, !Matched::AbstractChar)` where `T<:Number` at `char.jl:180`

...

1. `setindex! (::Vector{Float64}, ::Nothing, ::Int64)` @ `array.jl:839`
2. `top-level scope` @ `[Local: 9 [inlined]]`

```
• begin
•   # We initialize the random variables and preallocate storage
•   A = randn(Ny); As = zeros(Ny, T)
•   δ = rand(0:1, Ny); δs = zeros(Ny, T)
```

```

•   ε = rand(); es = zeros(T)
•   σ² = rand(); σ²s = zeros(T)
•   μ = randn(); μs = zeros(T)
•   for i in 1:T
•       μ = sample_μ(σ², θ, ν², y_outlier, δ, A, Ny); μs[i] = μ
•       σ² = sample_σ²(κ, Ny, λ, y_outlier, δ, A, μ); σ²s[i] = σ²
•       δ = sample_δ(ε, A, y_outlier, μ, σ²); δs[:, i] = δ
•       A = sample_A(τ², δ, y_outlier, μ, σ²); As[:, i] = A
•       ε = sample_ε(α, δ, β); es[i] = ε
•   end
• end

```

**UndefinedVarError: μs not defined**

1. top-level scope @ ( Local: 3

```

• begin
•   p1 = scatter(1:Ny, y_outlier, label = "y")
•   p2 = histogram(μs, label = "μ", normalize = true, lw = 0.0)
•   p3 = scatter(1:Ny, vec(mean(δs, dims = 2)), label = "δ")
•   p4 = histogram(es, label = "ε", normalize = true, lw = 0.0)
•   plot(p1, p2, p3, p4, legendfontsize=6.0)
• end

```

**(c) Which data points in the file *outlier.dat* are outliers? Use the samples generated in part (b) and the condition  $p(\delta_i|D) \geq 0.02$  in order to identify them.**

```

• begin
•   ## !! FILL IN !! Find for which i p(d_i|D) > 0.02
• end

```

fill\_in (generic function with 1 method)